

[illegible]

Student number	point total	req total	extra total	R1: joint positions (1p)	R2: joint rotations (2p)	R3: visualize joints (2p)	R4: skeletal SSD (4p)	R5: normal skinning (1p)	mod	notes / wtf / ...	GPU SSD (2p)	Wrist joint (5p)	IK (8p)	Animation (3p)	Dual quat (4p)	Other models (5p)	Pose capture (10p)	Style-based IK (10p)	Other extras (?p)	What other extras	
529992	19	10	9	1	2	2	4	1			2	4			3						
530185	0	0	0																		
530907	9.5	9.5	0	1	2	2	4	0.5		R5: Normals not normalized and treated as positions (w should be 0)											
530981	0	0	0																		
540094	0	0	0																		
540311	9.5	9.5	0	1	2	2	4	0.5		R5: normal transformed as a position and not normalized R3: You have used getRow when defining ahead when you should have used getCol; R5: Normals not normalized											
541543	9	9	0	1	2	1.5	4	0.5													
544375	10	10	0	1	2	2	4	1													
544566	0	0	0																		
549749	0	0	0																		
552969	0	0	0																		
556347	0	0	0																		
561578	0	0	0																		
563068	0	0	0																		
570116	0	0	0																		
586210	22	9.5	12.5	1	2	2	4	0.5		Your dual quaternion implementation is practically correct; GLSL treats matrices as arrays of column vectors stacked together, so the indices are transposed (column first, row second).	2	5			3.5				Dual quaternions on 2 the GPU		
587170	12	10	2	1	2	2	4	1			2										
587921	19.5	9	10.5	1	2	1.5	4	0.5		R3: Coordinate systems not scaled; R5: Normals not normalized; Dual quaternions: To get the shortest path of rotation, you should check the non-dual quaternions against each other and negate the weights for ones that point to a different direction.		5			3.5				Dual quaternions on 2 the GPU (2p)		
588137	0	0	0																		
588441	12	10	2	1	2	2	4	1			2										
589291	0	0	0																		
589437	9.5	9.5	0	1	2	2	4	0.5		R5: Incorrectly transforming with transpose matrix (also w should be 0). Normals not normalized.											
589848	0.5	0.5	0	0.5						R1 Visualization: You can get the joint position from the last column of the toWorld transformation matrix. R3: You should use transform matrix columns as the coordinate vectors, not rows; R5&GRU SSD: Normals not normalized. GPU SSD: you don't initialize some variables in the shader, which results in random data as initial value on newer nvidia gpus due to driver differences. Also, there appears one out of bounds error (i-4 should be i when reading from a.joints1). No points reduced but good to keep in mind.											
590112	10.5	8.5	2	1	2	1	4	0.5		R5: Normals not normalized and they are transformed as positions because getXYZ() is performed before multiplication. GPU SSD: Normals not normalized; Animation: It would look better if the if (nextPhase != 0) condition did not exist.	2			3							
590332	9.5	9.5	0	1	2	2	4	0.5													
590426	14.5	9.5	5	1	2	2	4	0.5			2										
593177	0	0	0																		
593452	9.5	9.5	0	1	2	2	4	0.5		R5: Incorrectly treating the normals as positions (w-coordinate 1). This causes the normals to have w ≠ 1 after transforming with the inverse transpose. ToCartesian() call then divides by the w-coordinate, giving incorrect results. Also, normals not normalized											
593876	9.5	9.5	0	1	2	2	4	0.5		R5: normals not normalized											
594367	11.5	9.5	2	1	2	2	4	0.5		R5 & GPU SSD: Normals not normalized	2										
594590	5	5	0	1	2	2															
594930	0	0	0																		
595201	15	9.5	5.5	1	2	2	4	0.5		R5&GPU SSD: Normals should have w = 0. GPU SSD: Normals not normalized. Wrist: Weights could have been better. Hand seems sort of disconnected and sticks out when rotated. Maybe the wrist is too high up the arm.	1.5	4									
595612	11.5	9.5	2	1	2	2	4	0.5		R5: normals not normalized (overly dark in tight creases)	2										
596048	11.5	9.5	2	1	2	2	4	0.5		R5: normals not normalized (overly dark in tight creases)	2										
596242	11.5	9.5	2	1	2	2	4	0.5		R5: normals not normalized (overly dark in tight creases)	2										
596747	0	0	0																		
596789	0	0	0																		
596792	9.5	9.5	0	1	2	2	4	0.5		R5: Normals not normalized											
596857	8.5	8.5	0	1	2	2	3.5	0		R4:The sum matrix is incorrectly initialized to identity instead of zero matrix; R5: Normals not transformed at all											
597445	15	10	5	1	2	2	4	1			2	1			2						
598088	12.5	10	2.5	1	2	2	4	1		GPU SSD: The vertex shader loops 8 elements even though the inputs are vec4 (this is a limitation of the vertex attribute API), you should loop 4 elements and do the same thing for a.joints2/aWeights2.	1.5	1									
598318	0	0	0																		
602851	0	0	0																		
603067	9.5	9.5	0	1	2	2	4	0.5		R5: normals not normalized (overly dark in tight creases)											
603096	5	5	0	1	2	2				R1: Instead of looping through all joints and using getJointParent(i) to see if i is a child, you could have just looped through j children.											
603326	0	0	0																		
604105	11	9.5	1.5	1	2	2	4	0.5		R5: normals not normalized (overly dark in tight creases). The vertex shader loops 8 elements even though the inputs are vec4 (this is a limitation of the vertex attribute API), you should loop 4 elements and do the same thing for a.joints2/aWeights2 -- could also be the reason why this didn't work everywhere. To improve perf, remove the inverse transposes -- it's a no-operation for a rigid transformation.	1.5										
606064	0	0	0																		
606268	0	0	0																		
608952	9	9	0	1	1.5	2	4	0.5		R2: m12 and m21 are switched when extracting the rotation matrix. R5: normal is transformed with an implicit w=1 and not normalized											
609142	21.5	9.5	12	1	2	2	4	0.5		R5&GPU SSD: normals not normalized; Wrist joint: Some unnatural stretching. Wrist joints are at least slightly in a wrong position. IK: Finite differences not ideal speedwise. Method looks solid, results leave some room for improvement.	2	3	7								
609155	0	0	0																		
609168	11	9.5	1.5	1	2	2	4	0.5		R5&GPU SSD: Normals not normalized; GPU SSD: You're missing the remaining 4 weights from your calculation.	1.5										
610827	4	4	0	1	2	1				R3: instead of adding scale to a specific axis, you should add right, up or ahead times scale to joint_world_pos and draw the second endpoint there.											
612155	0	0	0																		
612540	0	0	0																		
612812	0	0	0																		
621308	0	0	0																		
647175	0	0	0																		
647502	12	10	2	1	2	2	4	1			2										
648080	9	9.5	0	1	2	2	4	0.5	-0.5	Compiler errors (-0.5p); R5: normals should be transformed with 0 in w component since they are directions, not positions, and they are not normalized											
648569	10	10	0	1	2	2	4	1													
648660	0	0	0																		
649458	0	0	0																		
650191	0	0	0																		
650560	0	0	0																		
650829	9.5	9.5	0	1	2	2	4	0.5		R5: Normals not normalized											
651640	9.5	9.5	0	1	2	2	4	0.5		R5: normals not normalized (overly dark in tight creases)											
651802	0	0	0																		
652209	25	10	15	1	2	2	4	1		Animation: Use atan2, not atan, when converting quaternions back to euler angles. atan can only give values between -pi/2 and pi/2, meaning angles outside this range are calculated incorrectly. (It would have probably been more intuitive to have the animation last longer when frame count increases.)	2	5	2		4				Dual quaternions on 2 the GPU		
652584	25	10	15	1	2	2	4	1			2	5		2	4					Dual quaternions on 2 the GPU (2p)	
653156	0	0	0																		
653347	12	10	2	1	2	2	4	1		GPU SSD: you don't initialize your vectors in the shader, which results in random data as initial value on newer nvidia gpus due to driver differences. No points reduced but good to keep in mind.	2										
654142	9.5	9.5	0	1	2	2	4	0.5		R5: Normal is not a point but a direction, i.e. w should be 0, and normals should be normalized											
654294	9.5	9.5	0	1	2	2	4	0.5		R5: Normals not normalized											

Student number	point total	req total	extra total	R1: joint positions (1p)	R2: joint rotations (2p)	R3: visualize joints (2p)	R4: skeletal SSD (4p)	R5: normal skinning (1p)	mod	notes / wtf / ...	GPU SSD (2p)	Wrist joint (5p)	IK (8p)	Animation (3p)	Dual quat (4p)	Other models (5p)	Pose capture (10p)	Style-based IK (10p)	Other extras (7p)	What other extras	
654618	11.5	9.5	2		1	2	2	4	0.5	R5: normals not normalized (overly dark in tight creases). Normal w-component should also be 0 in principle; here the inverse transpose effectively sets it to 0 anyway.	2										
655109	0	0	0																		
655361	0	0	0																		
655390	0	0	0																		
656014	8	8	0	1	1.5	1	4	0.5		R2: Rotations were supposed to be calculated in setJointRotation; R3: You forgot to draw the bones; R5: Normals not normalized											
657068	8.5	8.5	0	1	2	2	3.5			R4: Mat4f initializes to identity by default. Since you're gathering a sum, you need to zero initialize the matrix first.											
657181	11.5	9.5	2	1	2	2	4	0.5		R5 & GPU SSD: Normals not normalized	2										
657437	9.5	9.5	0	1	2	2	4	0.5		R5: Normals transformed as positions and not normalized											
										R5&GPU SSD: Normals not normalized and treated as positions (w should be 0); Dual quaternions: To get the shortest path of rotation, you should check the non-dual quaternions against each other and negate the weights for ones that point to a different direction. IK: Ok attempt, the key parts seem to be there.	2	5	4	3	3.5				2	Dual quaternions on the GPU (2p)	
657767	29	9.5	19.5	1	2	2	4	0.5													
657893	0	0	0																		
663434	0	0	0																		
665173	9.5	9.5	0	1	2	2	4	0.5		R5: normals not normalized (overly dark in tight creases)											
										R5&GPU SSD: normals transformed as positions and not normalized; Dual quaternions: To get the shortest path of rotation, you should check the non-dual quaternions against each other and negate the weights for ones that point to a different direction.	2	5		3	3.5	1		2	Dual quaternions on the GPU (2p)		
665678	26	9.5	16.5	1	2	2	4	0.5													
666208	0	0	0																		
666211	10	10	0	1	2	2	4	1													
666253	0	0	0																		
710015	11.5	9.5	2	1	2	2	4	0.5		R5 & GPU SSD: Normals not normalized	2										
715298	12	10	2	1	2	2	4	1			2										
716734	0	0	0																		
717377	0	0	0																		
717539	0	0	0																		
718020	0	0	0																		
										R5: normals not normalized (overly dark in tight creases). Nice animation system, but for a more practical solution you might want to consider only storing the keyframes and interpolating on the fly.	2			3							
718208	14.5	9.5	5	1	2	2	4	0.5		R5: normals should be normalized after the sum, not per every term.											
718512	9.5	9.5	0	1	2	2	4	0.5		R5: normals not normalized (overly dark in tight creases)	2	5		3							
718826	19.5	9.5	10	1	2	2	4	0.5		Submission zip doesn't open with standard Windows zip tool, 7zip opens it but it's almost empty.											
719032	0	0	0																		
721619	0	0	0																		
721923	0	0	0																		
723154	0	0	0																		
										R2: On row 43 you are have Rx when you should have Rz. This causes the model to look flat when rotating around the x-axis. It would have been simpler to use the given Mat3f.rotation() function; R5: Normals transformed as positions. Your solution for GPU SSD is correct though.											
723329	10.5	8.5	2	1	1	2	4	0.5		R3: The columns of the transformation matrix would have given you the correct directions for the coordinate systems. Now the lines always point in the same direction. R5: Normals not normalized; Wrist: No weights and because left_hand is added in the middle of the file, the weights are off for some joints; GPU SSD: you don't initialize your vectors in the shader, which results in random data as initial value on newer nvidia gpus due to driver differences. No points reduced but good to keep in mind.	2										
723468	11.5	8.5	3	1	2	1	4	0.5		R5: Normals not normalized; GPU SSD: you don't initialize your matrix in the shader, which results in random data as initial value on newer nvidia gpus due to driver differences. No points reduced but good to keep in mind; Wrist: No weights and because left_wrist is added in the middle of the file, the weights are off for some joints. Wrists are not in the correct location.	2	1									
										R5: Normals not normalized; GPU SSD: The ssd code is exactly right apart from the second loop bounds. After the ssd, you should use the new normal for calculating clampedCosine and the new position for calculating gl_Position.	1										
726915	10.5	9.5	1	1	2	2	4	0.5													
728696	0	0	0																		
729297	0	0	0																		
732323	0	0	0																		
737551	12	10	2	1	2	2	4	1			2										
765714	0	0	0																		
765756	0	0	0																		
										Dual quaternions: To get the shortest path of rotation, you should check the real parts of the quaternions against each other and negate the weights for ones that point to a different direction. R4: computeToBindTransforms and getSSDTransforms incorrect; R5: This should just be a matrix vector multiplication just as you did for positions.	2	5			3.5						
765785	20.5	10	10.5	1	2	2	4	1													
765882	7	7	0	1	2	2	2	0													
766108	0	0	0																		
767136	0	0	0																		
769396	0	0	0																		
										In R2 you're setting the wrong elements from to_parent; should be m3() instead of m()i3 (or just zeros). This makes the model completely distorted under any joint movement. In R3, the last coordinates of right and up are read from the wrong columns; are 1 and 2 but should be 0 and 1.											
772419	8.5	8.5	0	1	1	1.5	4	1													
784465	0	0	0																		
784847	0	0	0																		
784902	0	0	0																		
785053	0	0	0																		
785134	9.5	9.5	0	1	2	2	4	0.5		R5: normals not normalized (sometimes overly dark in tight creases)											
785163	0	0	0																		
785228	11.5	9.5	2	1	2	2	4	0.5		R5: normals not normalized (overly dark in tight creases). Good point about the key bindings, we'll do something about this.	2										
										R5: Don't apply translation to normals, i.e. set the w coordinate to 0 before transforming. The transposed() call doesn't do anything right now, use transpose() to modify the underlying object. Only transposing would give incorrect results, you need the inverse transpose. However in this case the transformation matrix is orthogonal which means it's inverse equals it's transpose and therefore the inverse transpose is just the original matrix. Also, normals not normalized; GPU SSD: Normals not transformed											
785257	11	9.5	1.5	1	2	2	4	0.5			1.5										
785325	0	0	0																		
785354	9.5	9.5	0	1	2	2	4	0.5		R5: Normals not normalized											
785367	9.5	9.5	0	1	2	2	4	0.5		R5: Normals not normalized											
785435	0	0	0																		
785448	0	0	0							Submitted but nothing done											
785451	5	5	0	1	2	2															
										README.txt empty. R1: incorrect order of transformations (in the end we multiply with the position from the right; we go first into the parent bone's state and from there to the world). R3: the idea is to visualize the local coordinate system of the bone; instead of adding scale to a specific axis, you should add right, up or ahead times scale to joint_world_pos. R4: not zeroing the position between loops (you just keep on summing), unfortunate order of computations (do the matrix product in parentheses to get the correct answer.)											
785493	6	6.5	0	0.5	2	1	3	-0.5													
785503	0	0	0																		
785516	0	0	0																		
795551	10	10	0	1	2	2	4	1													
795577	10	10	0	1	2	2	4	1													
795593	5	5	0	1	2	2															
795629	0	0	0																		
795658	0	0	0																		
795674	12	9.5	2.5	1	2	2	4	0.5		R5&GPU SSD: Normals not normalized; GPU SSD: Normals should have w = 0	1.5				1						

[illegible]