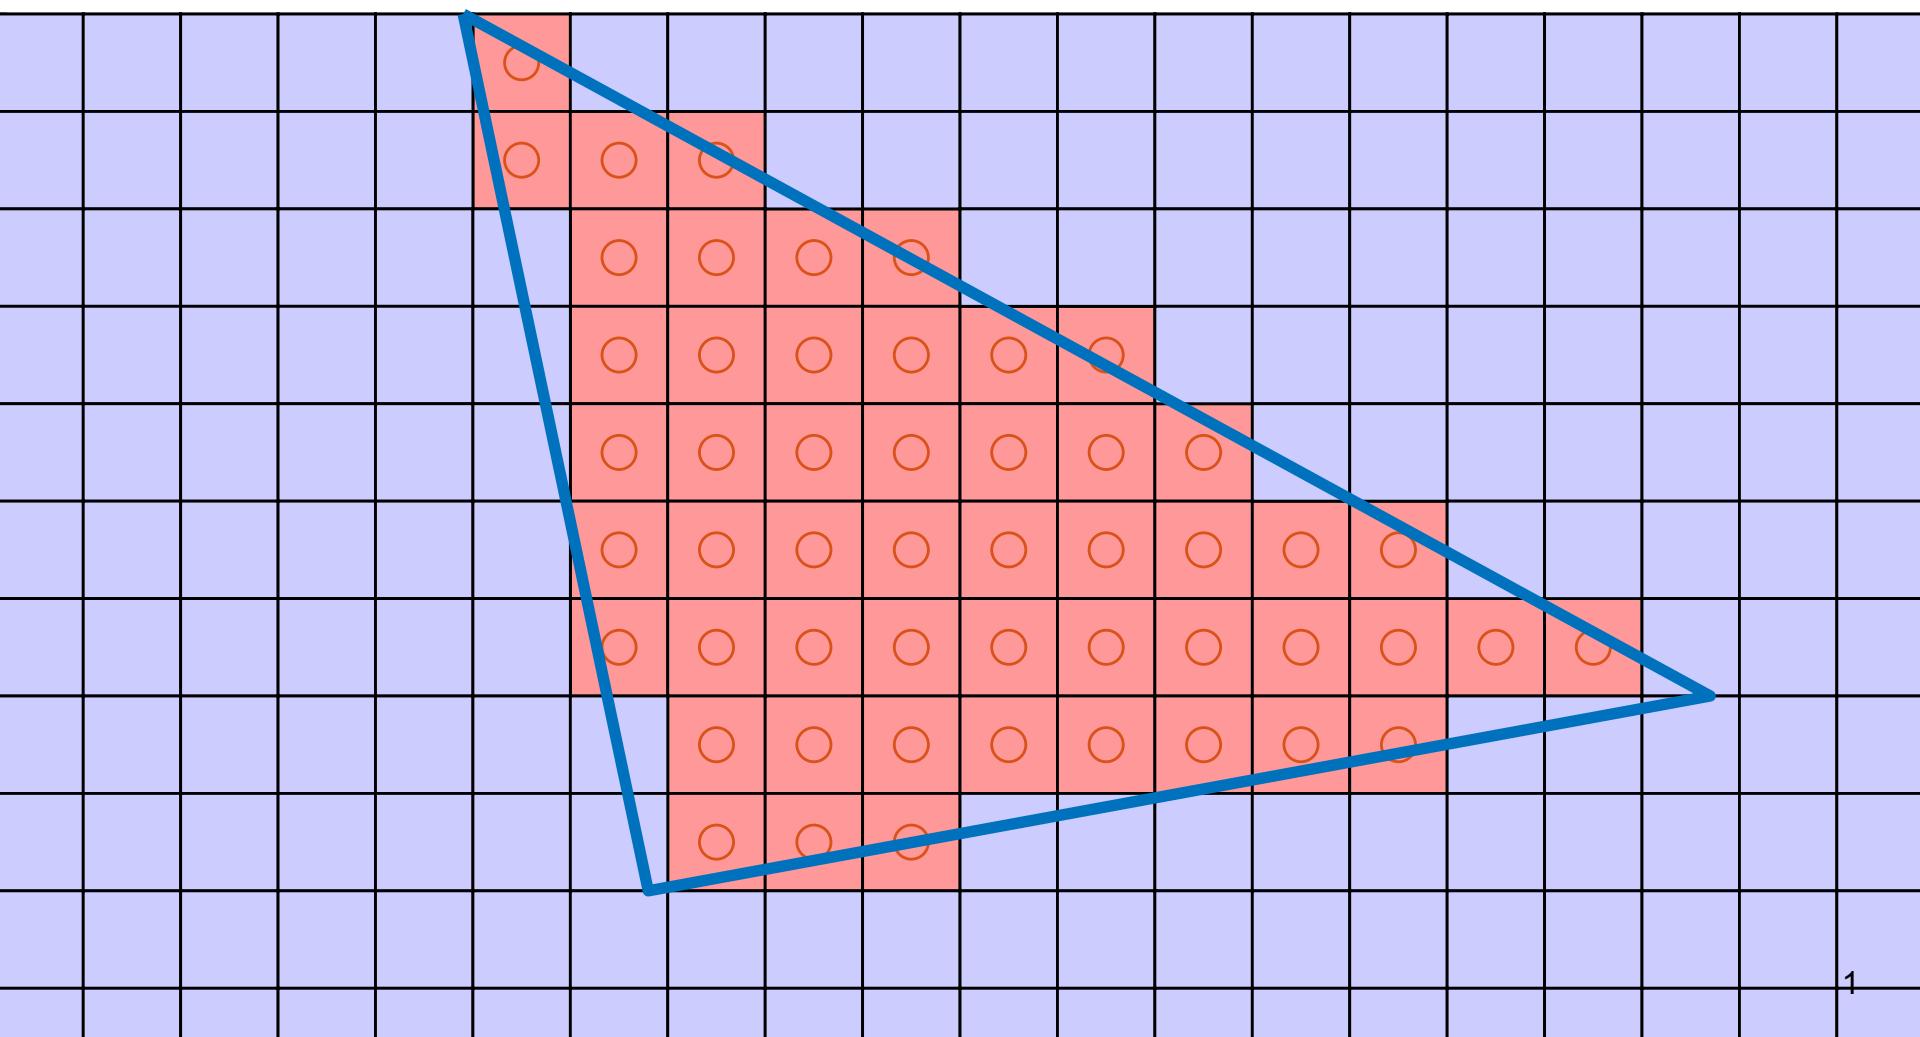


Rasterization & The Graphics Pipeline

15.1 Basics: Rasterization vs. Ray Tracing

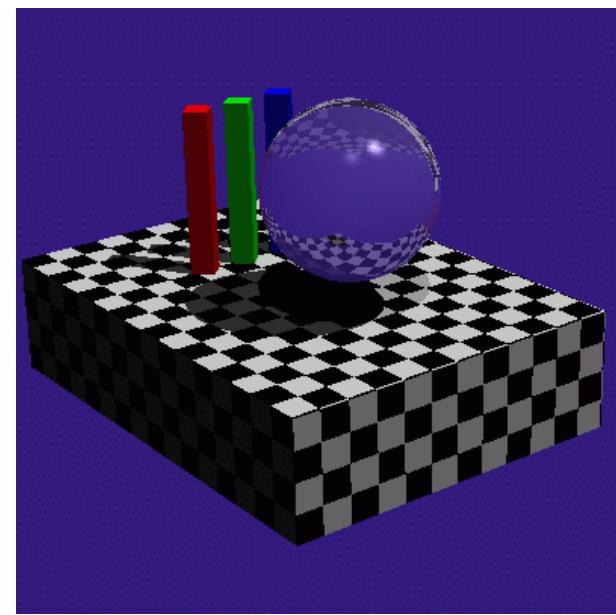
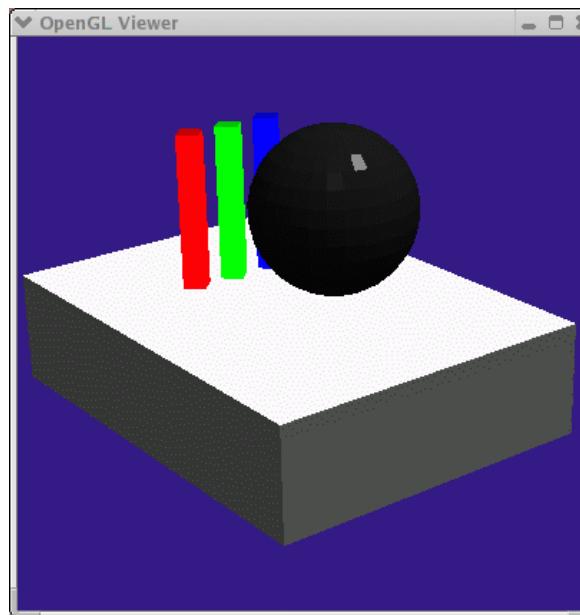
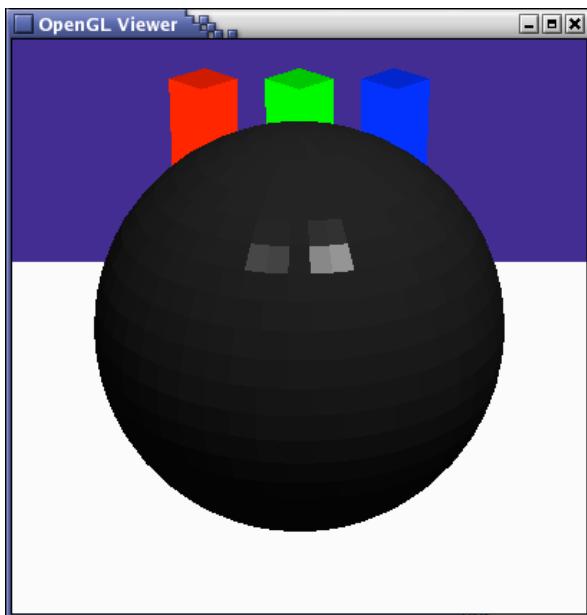


In This Video

- What is real-time graphics mostly based on?
- Rasterization vs. ray tracing
 - Different ordering of pixel/primitive combinations
 - Working sets of the two algorithms
 - What is streamed over and then forgotten?
 - In contrast, what is kept in memory for random access?
 - Advantages and disadvantages

How Do We Render Interactively?

- Use graphics hardware, via OpenGL or DIRECTX
 - OpenGL is multi-platform, DirectX is MS only

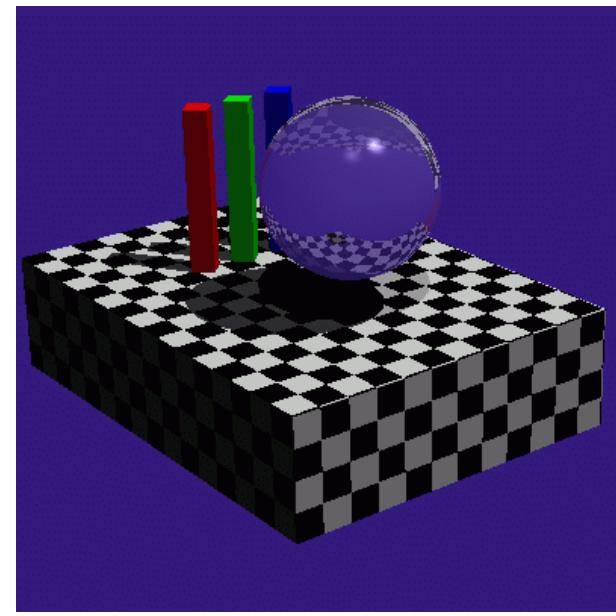
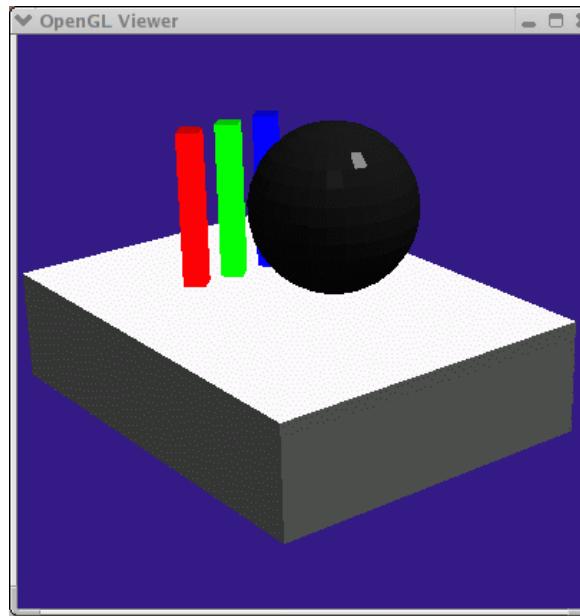
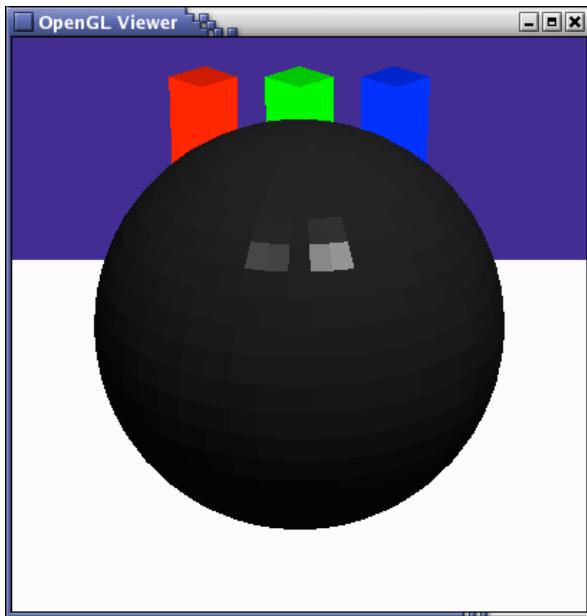


- *OpenGL rendering* Very basic!

Our ray tracer

How Do We Render Interactively?

- Use graphics hardware, via OpenGL or DIRECTX
 - OpenGL is multi-platform, DirectX is MS only



OpenGL rendering

Our ray tracer

- Most global effects available in ray tracing will be sacrificed for speed, but some can be approximated

All Is Not Lost



- Example: “Screen space reflection” is possible..
- ... but MUCH more cumbersome than ray tracing

Uncharted 4



Quantum Break



Battlefield 1 by DICE



Ray Casting vs. GPUs for Triangles

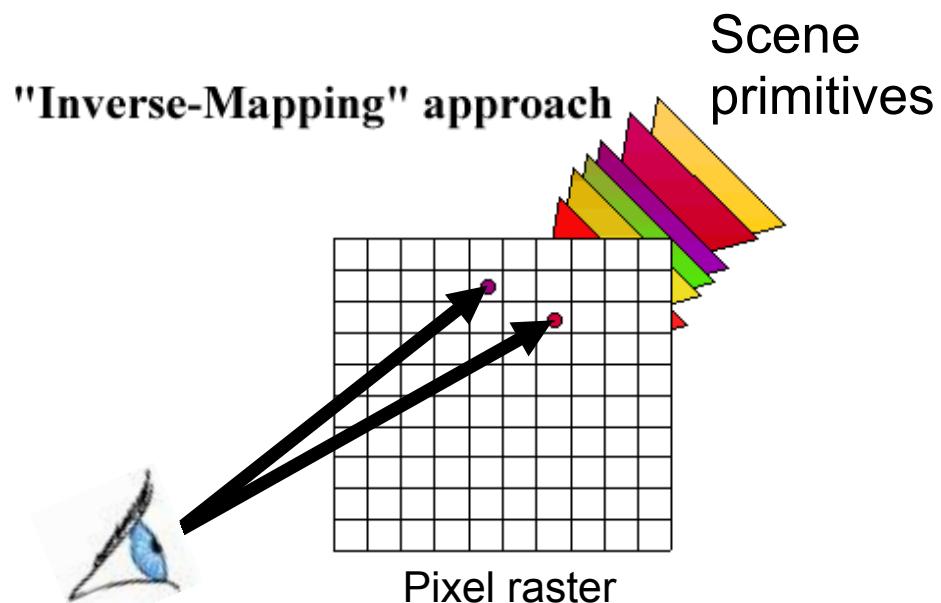
Ray Casting

For each pixel (ray)

 For each triangle

 Does ray hit triangle?

 Keep closest hit



Ray Casting vs. GPUs for Triangles

Ray Casting

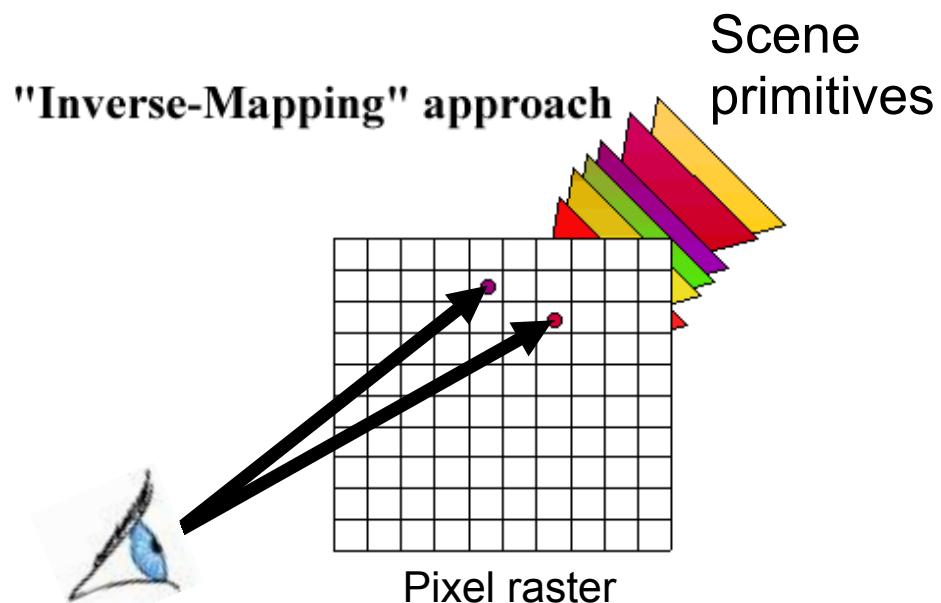
For each pixel (ray)

For each triangle

Does ray hit triangle?

Keep closest hit

Last time we replaced this with a traversal of a Bounding Volume Hierarchy (BVH)



Ray Casting vs. GPUs for Triangles

Ray Casting

For each pixel (ray)

For each triangle

Does ray hit triangle?

Keep closest hit

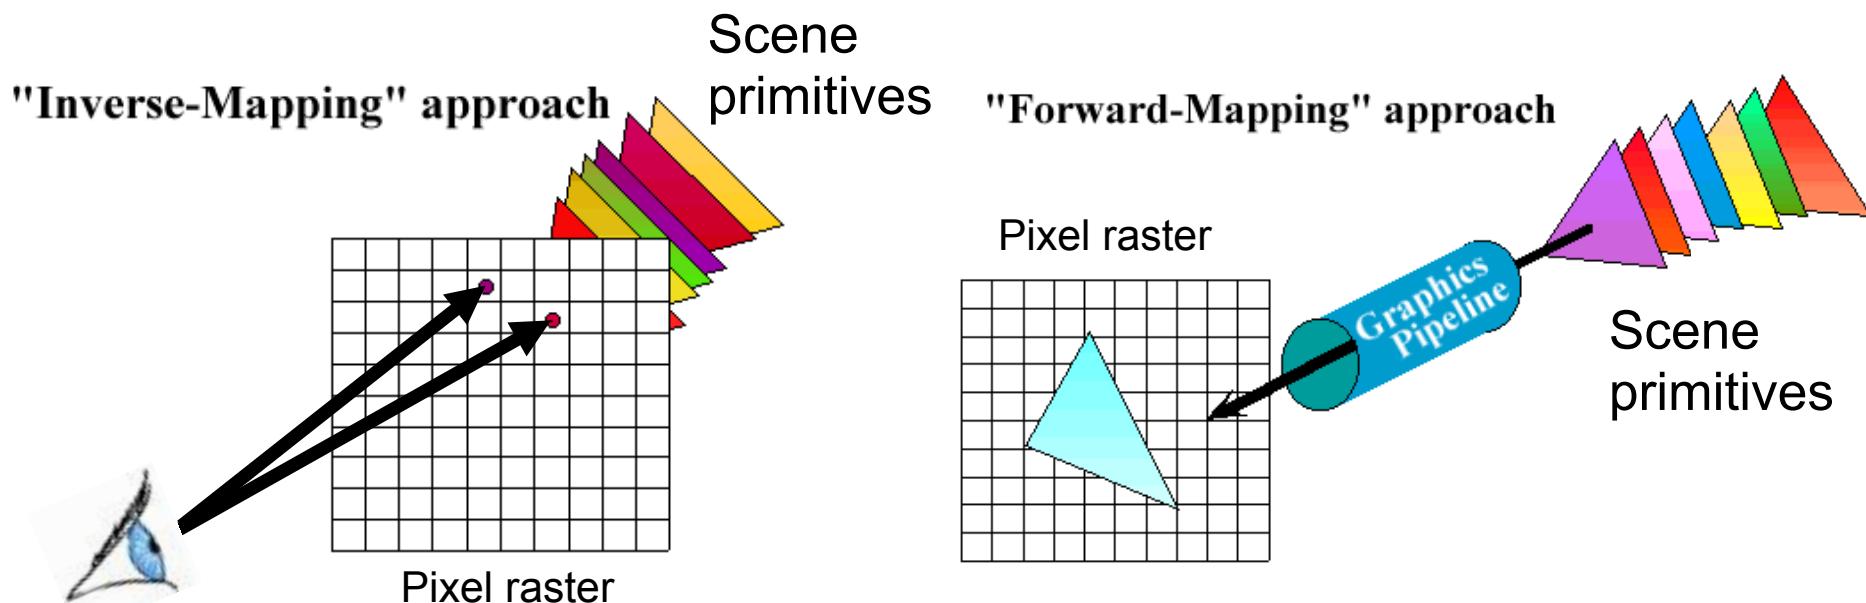
GPU

For each triangle

For each pixel

Does triangle cover pixel?

Keep closest hit



Ray Casting vs. GPUs for Triangles

Ray Casting

For each pixel (ray)

For each triangle

Does ray hit triangle?

Keep closest hit

GPU

For each triangle

For each pixel

Does triangle cover pixel?

Keep closest hit

It's just a different order of the loops!

GPUs do Rasterization

- The process of taking a triangle and figuring out which pixels it covers is called **rasterization**

GPU

For each triangle

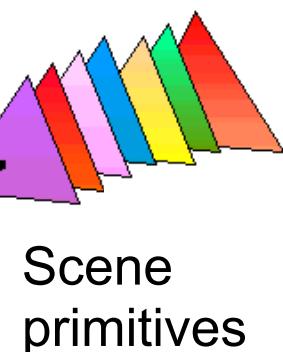
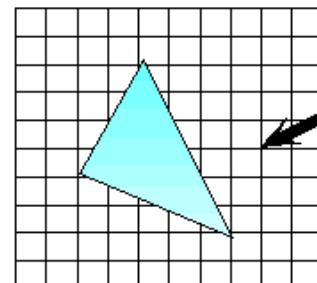
For each pixel

Does triangle cover pixel?

Keep closest hit

"Forward-Mapping" approach

Pixel raster



Scene
primitives

GPUs do Rasterization

- The process of taking a triangle and figuring out which pixels it covers is called **rasterization**
- We've seen acceleration structures for ray tracing; rasterization is not stupid either
 - We're not actually going to test *all* pixels for each triangle

GPU

For each triangle

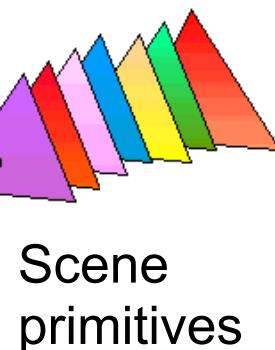
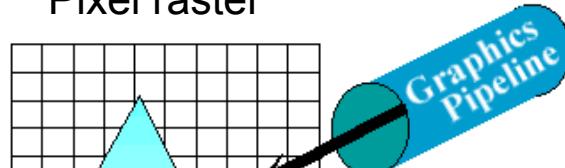
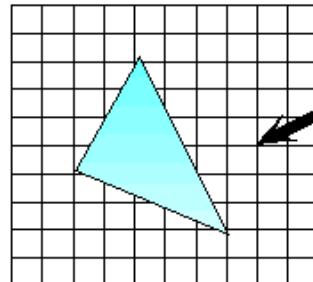
For each pixel

Does triangle cover pixel?

Keep closest hit

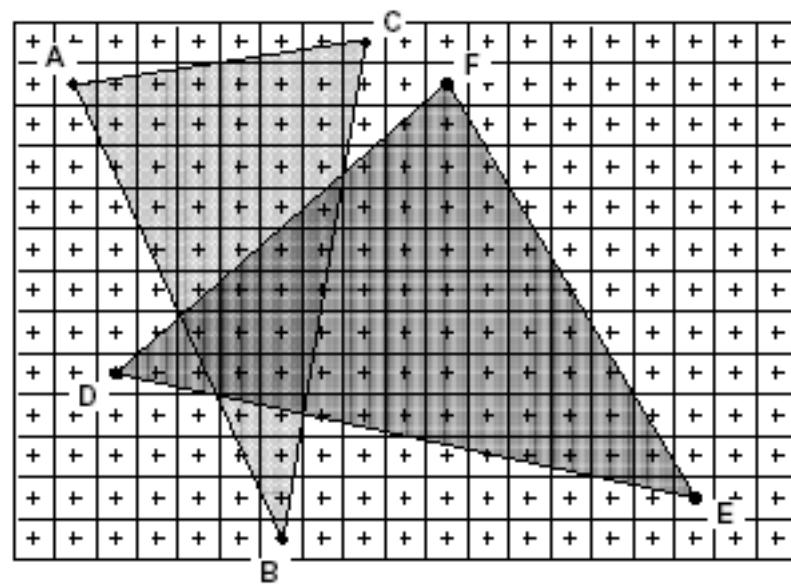
"Forward-Mapping" approach

Pixel raster



Rasterization (“Scan Conversion”)

- Given a triangle’s vertices & extra info for shading, figure out which pixels to "turn on" to render the primitive
- Compute illumination values to "fill in" the pixels within the primitive
- At each pixel, keep track of the closest primitive (z-buffer)
 - Only overwrite if triangle being drawn is closer than the previous triangle in that pixel



What are the Main Differences?

Ray Casting

For each pixel (ray)

For each triangle

Does ray hit triangle?

Keep closest hit

GPU

For each triangle

For each pixel

Does triangle cover pixel

Keep closest hit

Ray-centric

Triangle-centric

- What needs to be stored in memory in each case?

What are the Main Differences?

Ray Casting

For each pixel (ray)

For each triangle

Does ray hit triangle?

Keep closest hit

GPU

For each triangle

For each pixel

Does triangle cover pixel

Keep closest hit

Ray-centric

Triangle-centric

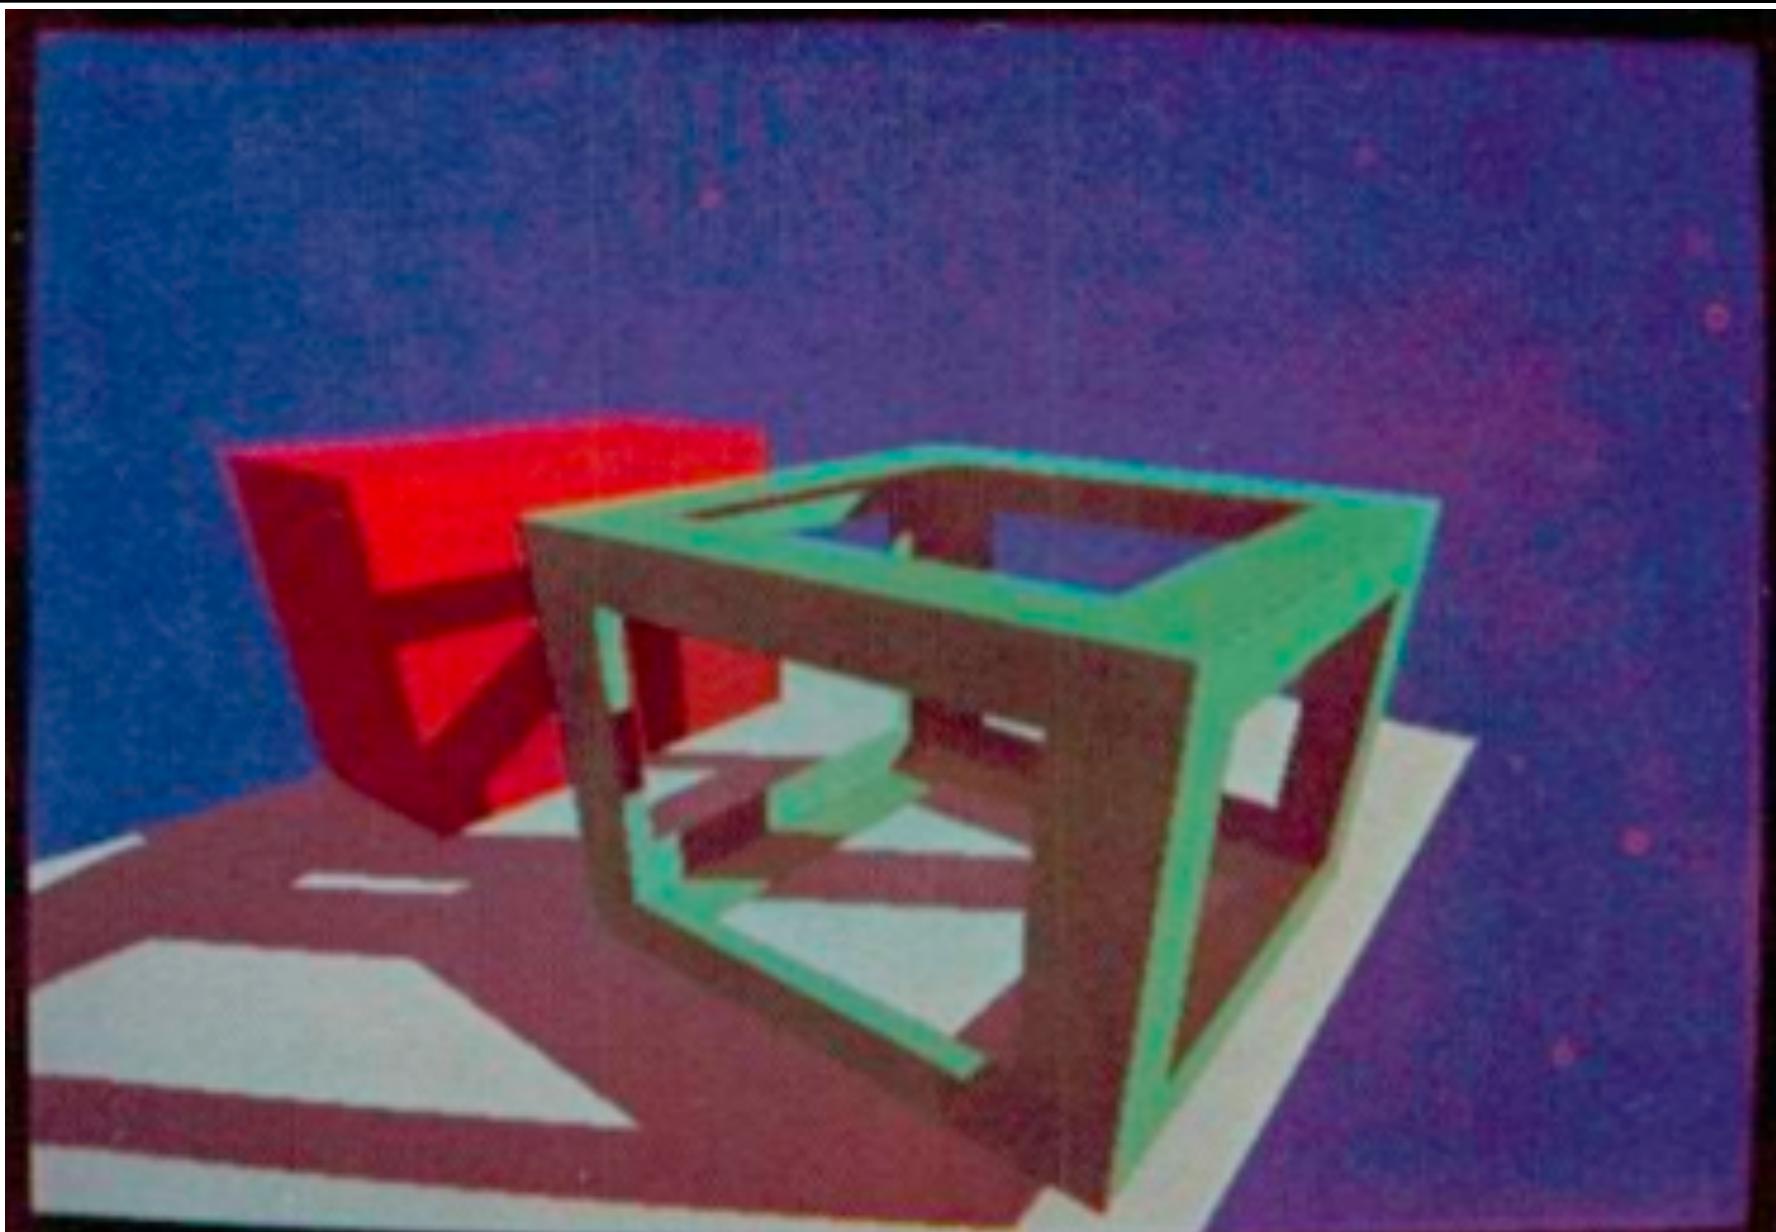
- In this basic form, ray tracing needs the entire scene description in memory at once
 - Then, can sample the image completely freely
- The rasterizer only needs one triangle at a time, *plus* the entire image and associated depth information for all pixels

Rasterization Advantages

- Modern scenes are more complicated than images
 - A 1920x1080 frame at 64-bit color and 32-bit depth per pixel is 24MB (not that much)
 - Of course, if we have more than one sample per pixel this gets larger, but e.g. 4x supersampling is still a relatively comfortable ~100MB
 - *4K gets to hundreds of megabytes!*
 - Our scenes are routinely larger than this in bytes
 - This wasn't always true...

Rasterization Advantages

Weiler, Atherton 1977

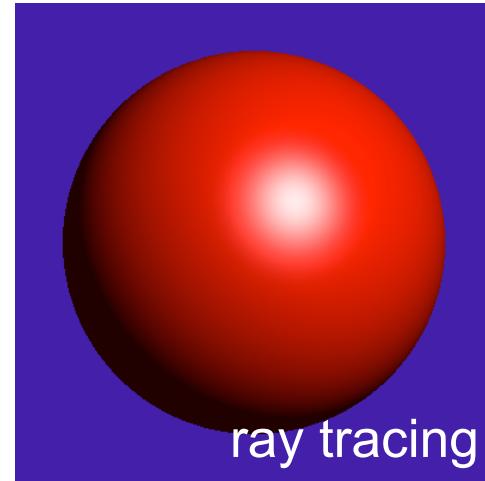


Rasterization Advantages

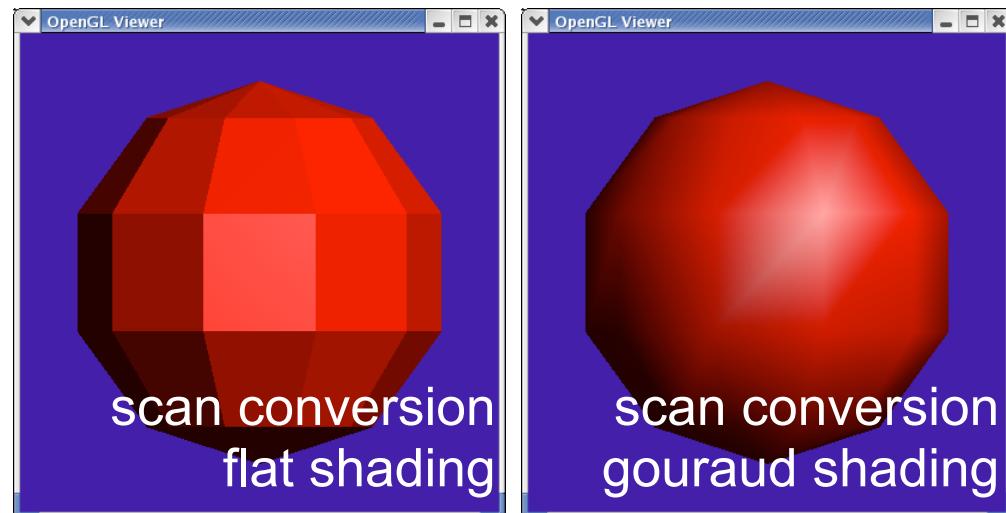
- Modern scenes are more complicated than images
 - A 1920x1080 frame (1080p) at 64-bit color and 32-bit depth per pixel is 24MB (not that much)
 - Of course, if we have more than one sample per pixel (later) this gets larger, but e.g. 4x supersampling is still a relatively comfortable ~100MB
 - Our scenes are routinely larger than this
 - This wasn't always true
- A rasterization-based renderer can *stream* over the triangles, no need to keep entire dataset around
 - Allows parallelism and optimization of memory systems

Rasterization Limitations

- Restricted to scan-convertible primitives
 - Pretty much: triangles
- Faceting, shading artifacts
 - *This has largely gone away with programmable per-pixel shading, though*
- No unified handling of shadows, reflection, transparency
- Potential problem of overdraw (high depth complexity)
 - Each pixel touched many times



ray tracing



scan conversion
flat shading

scan conversion
gouraud shading

Ray Casting / Tracing



- Advantages
 - Generality: can render anything that can be intersected with a ray
 - Easily allows recursion (shadows, reflections, etc.)

- Disadvantages
 - Harder to implement in hardware (less computation coherence, must fit entire scene in memory, worse memory behavior)
 - Not such a big point any more given general purpose GPUs
 - Was traditionally too slow for interactive applications
 - **These are not as true any more as they used to be**

Up next: the Graphics Pipeline



BATTLEFIELD 4

Battlefield 4 by DICE (Stockholm)