


Rasterization & The Graphics Pipeline

15.3 Rasterization using edge functions





In This Video

- Edge functions: finding out which pixels are covered
 - and various optimisations
- Extra: Projection matrices
 - From 3D to 2D via homogeneous coordinates


The Graphics Pipeline

- Project vertices to 2D (image)
 - We now have screen coordinates
- Rasterize triangle: find which pixels should be lit
- Compute per-pixel color
- Test visibility (Z-buffer), update frame buffer




2D Scan Conversion

- Primitives are “continuous” geometric objects; screen is discrete (pixels)




2D Scan Conversion

- Primitives are “continuous” geometric objects; screen is discrete (pixels)
- Rasterization computes a discrete approximation in terms of pixels (**how?**)




Edge Functions




Edge Functions

- The triangle's 3D edges project to line segments in the image (thanks to planar perspective!)
 - Lines map to lines, not curves




Edge Functions

- The triangle's 3D edges project to line segments in the image (thanks to planar perspective)




Edge Functions

- The triangle's 3D edges project to line segments in the image (thanks to planar perspective)
- The interior of the triangle is the set of points that is inside all three halfspaces defined by these lines



Edge Functions

- The triangle's 3D edges project to line segments in the image (thanks to planar perspective)
- The interior of the triangle is the set of points that is inside all three halfspaces defined by these lines




$$E_i(x, y) = a_i x + b_i y + c_i$$

$$(x, y) \text{ within triangle} \\ \Leftrightarrow \\ E_i(x, y) \geq 0, \\ \forall i = 1, 2, 3$$

Edge Function 1

$$E_1(x, y) < 0$$

$$E_1(x, y) > 0$$



$$E_i(x, y) = a_i x + b_i y + c_i$$

(x, y) within triangle


\Leftrightarrow

$$E_i(x, y) \geq 0,$$

$$\forall i = 1, 2, 3$$

Edge Function 2

$$E_2(x, y) < 0$$



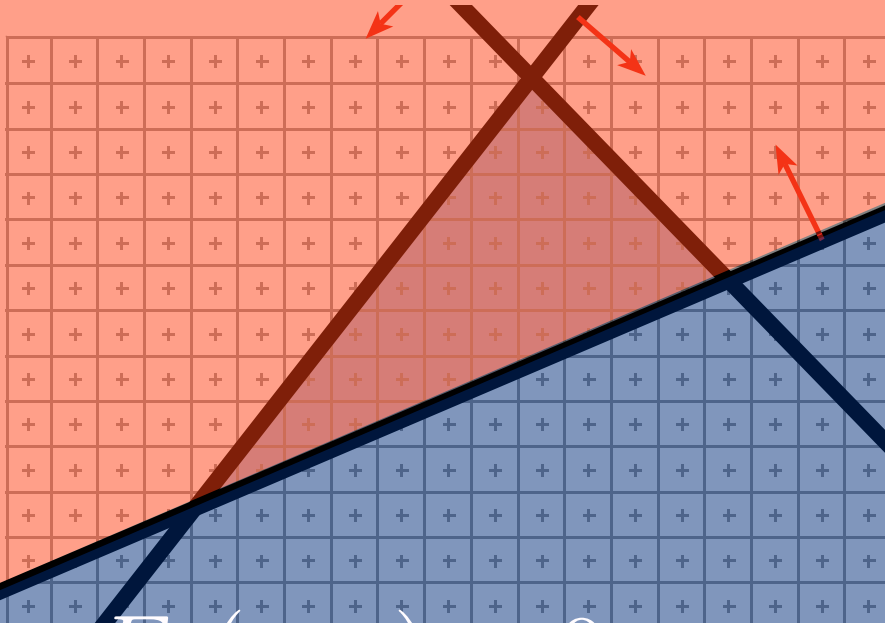
$$E_i(x, y) = a_i x + b_i y + c_i$$

(x, y) within triangle

$$E_i(x, y) \geq 0, \\ \forall i = 1, 2, 3$$

Edge Function 3

$$E_3(x, y) > 0$$




$$E_3(x, y) < 0$$

$$E_i(x, y) = a_i x + b_i y + c_i$$

(x, y) within triangle


$$E_i(x, y) \geq 0, \quad \forall i = 1, 2, 3$$

Brute Force Rasterizer




Brute Force Rasterizer

- Compute E_1 - E_3 from projected vertices
 - Called “triangle setup”, yields a_i, b_i, c_i for $i=1,2,3$




Brute Force Rasterizer

- Compute E_1 - E_3 from projected vertices
 - Called “triangle setup”, yields a_i, b_i, c_i for $i=1,2,3$
- For each pixel (x, y)
 - Evaluate edge functions at pixel center
 - If all non-negative, pixel is in!



Brute Force Rasterizer


- Compute E_1 - E_3 from projected vertices
 - Called “triangle setup”, yields a_i, b_i, c_i for $i=1,2,3$
- For each pixel (x, y)
 - Evaluate edge functions at pixel center
 - If all non-negative, pixel is in!



Problem?


Brute Force Rasterizer

- Compute E_1 - E_3 from projected vertices (“setup”)
- For each pixel (x, y)
 - Evaluate edge functions at pixel center
 - If all non-negative, pixel is in!




If the triangle is small, lots of useless computation if we really test all pixels

Easy Optimization




Easy Optimization

- Improvement: Scan over only the pixels that overlap the *screen bounding box* of the triangle




Easy Optimization

- Improvement: Scan over only the pixels that overlap the *screen bounding box* of the triangle
- How do we get such a bounding box?



Easy Optimization

- Improvement: Scan over only the pixels that overlap the *screen bounding box* of the triangle
- How do we get such a bounding box?
 - X_{\min} , X_{\max} , Y_{\min} , Y_{\max} of the projected triangle vertices



Rasterization Pseudocode

**Note: No
visibility**

For every triangle

 Compute projection for vertices, compute the E_i


 Compute bbox, clip bbox to screen limits

 For all pixels in bbox

 Evaluate edge functions E_i

 If all > 0

 Framebuffer[x,y] = triangleColor



**Bounding box clipping is easy,
just clamp the coordinates to
the screen rectangle**

Can We Do Better?

For every triangle

 Compute projection for vertices, compute the E_i


 Compute bbox, clip bbox to screen limits

 For all pixels in bbox

Evaluate edge functions $a_i x + b_i y + c_i$

 If all > 0

 Framebuffer[x,y] = triangleColor



Can We Do Better?

For every triangle

 Compute projection for vertices, compute the E_i


 Compute bbox, clip bbox to screen limits

 For all pixels in bbox

Evaluate edge functions $a_i x + b_i y + c_i$

 If all > 0

 Framebuffer[x,y] = triangleColor



These are linear functions of the pixel coordinates (x,y), i.e., they only change by a constant amount when we step from x to x+1 (resp. y to y+1)

Incremental Edge Functions


```
For every triangle
  ComputeProjection
  Compute bbox, clip bbox to screen limits
  For all scanlines y in bbox
    Evaluate all  $E_i$ 's at  $(x_0, y)$ :  $E_i = a_i x_0 + b_i y + c_i$ 
    For all pixels x in bbox
      If all  $E_i > 0$ 
        Framebuffer[x,y ] = triangleColor
    Increment line equations:  $E_i += a_i$ 
```

- We save ~two multiplications and two additions per pixel when the triangle is large

Incremental Edge Functions

```
For every triangle
  ComputeProjection
  Compute bbox, clip bbox to screen limits
  For all scanlines y in bbox
    Evaluate all  $E_i$ 's at  $(x_0, y)$ :  $E_i = a_i x_0 + b_i y + c_i$ 
    For all pixels x in bbox
      If all  $E_i > 0$ 
        Framebuffer[x, y] = triangleColor
    Increment line equations:  $E_i += a_i$ 
```


- We save ~two multiplications and two additions per pixel when the triangle is large



Can also zig-zag to avoid reinitialization per scanline, just initialize once at x_0, y_0


Questions?

- For a really HC piece of rasterizer engineering, see the hierarchical Hilbert curve rasterizer by McCool, Wales and Moule.
 - (Hierarchical? We'll look at that next..)




Can We Do Even Better?

- We compute the line equation for many useless pixels
- What could we do?




Indeed, We Can Be Smarter



Indeed, We Can Be Smarter


- Hierarchical rasterization!
 - Conservatively test **blocks of pixels** before going to per-pixel level (can skip large blocks at once)
 - Usually two levels



Conservative tests of axis-aligned blocks vs. edge functions are not very hard, thanks to linearity. See Akenine-Möller and Aila, Journal of Graphics Tools 10(3), 2005.

Indeed, We Can Be Smarter


- Hierarchical rasterization!
 - Conservatively test **blocks of pixels** before going to per-pixel level (can skip large blocks at once)
 - Usually two levels



Can also test if an entire block is **inside** the triangle; then, can skip edge functions tests for all pixels for even further speedups. (Must still test Z, because they might still be occluded.)


Oldskool Rasterization

- Compute the boundary pixels using line rasterization




Oldskool Rasterization

- Compute the boundary pixels using line rasterization
- Fill the spans



Oldskool Rasterization

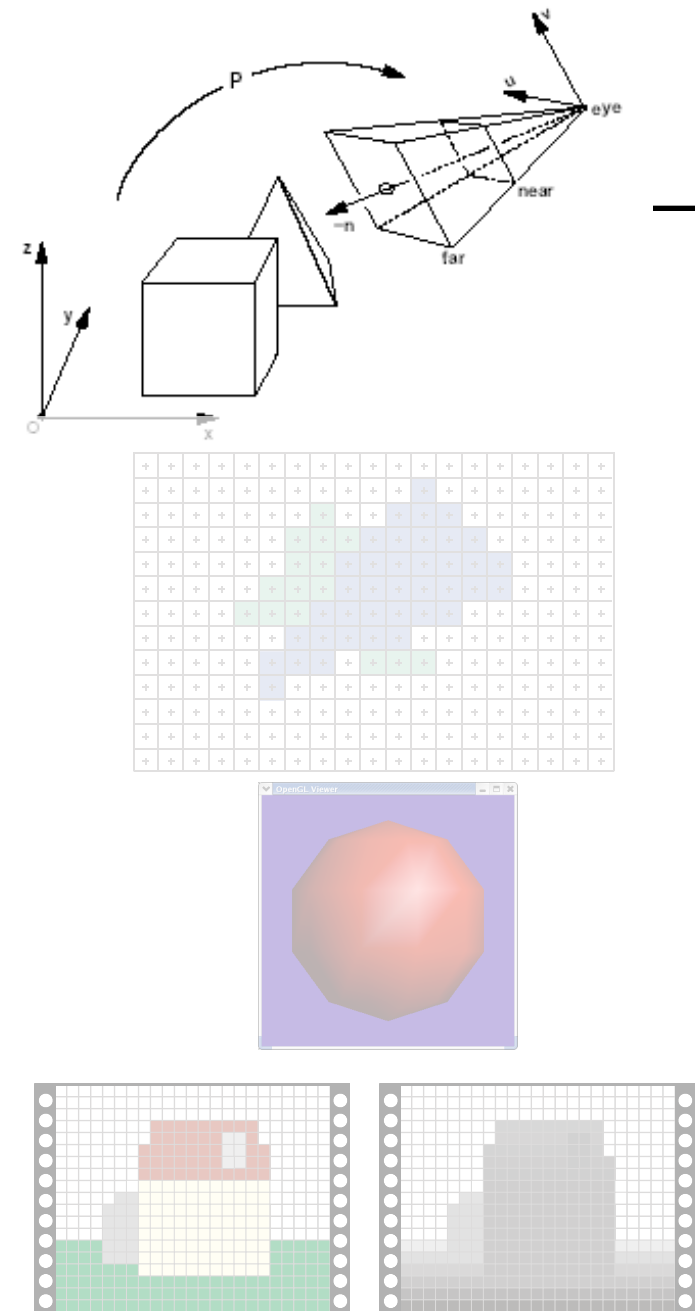
- Compute the boundary pixels using line rasterization
- Fill the spans



MUCH MORE annoying
to implement than edge
functions


Projection

- Project vertices to 2D (image)
- Rasterize triangle: find which pixels should be lit
- Compute per-pixel color
- Test visibility (Z-buffer), update frame buffer




Orthographic vs. Perspective

- Orthographic




- Perspective




The View Frustum in 2D

- (In 3D this is a truncated pyramid.)




The View Frustum in 2D

- We can transform the frustum by a modified projection in a way that makes it a square (cube in 3D) after division by w' .



The View Frustum in 2D


- We can transform the frustum by a modified projection in a way that makes it a square (cube in 3D) after division by w' .



The final image is obtained by merely dropping the z coordinate after projection (orthogonal projection)


The View Frustum in 2D

- We can transform the frustum by a modified projection in a way that makes it a square (cube in 3D) after division by w' .




- The x' coordinate does not change w.r.t. the usual flattening projection, i.e., x'/w' stays the same

The Canonical View Volume



- Point of the exercise: This gives screen coordinates and depth values for Z-buffering with unified math
 - Caveat: OpenGL and DirectX define Z differently $[0,1]$ vs. $[-1,1]$

OpenGL Form of the Projection

$$\begin{pmatrix} x' \\ y' \\ z' \\ w' \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & \frac{\text{far} + \text{near}}{\text{far} - \text{near}} & -\frac{2 * \text{far} * \text{near}}{\text{far} - \text{near}} \\ 0 & 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix}$$



**Homogeneous coordinates
within canonical view volume**

**Input point in
view coordinates**

- Details/more intuition in handout in MyCourses
 - “Understanding Projections and Homogenous Coordinates”

Recap: Projection

- Perform rotation/translation/other transforms to put viewpoint at origin and view direction along z axis
- Combine with projection matrix (perspective or orthographic)
 - Homogenization achieves foreshortening
- **Corollary:** The entire transform from object space to canonical view volume $[-1,1]^3$ is a single matrix



© www.scratchapixel.com